

# Monitor Object

## Concurrency Pattern

Huib van den Brink

Center for Software Technology, Utrecht University  
<http://www.cs.uu.nl/groups/ST/>

October 14, 2005



# Outline

- 1 Concurrency problems
  - Problems involved
  - Active Object
- 2 Monitor Object
  - The pattern
  - Metaphor
  - Synchronized implementation example
  - Lock on parts of method / among several objects
  - Introduced problems
  - Performance issues
  - Pros & Cons
- 3 Finally
  - Questions
  - Exercise



# Computers aren't woman

## Concurrency problems

- ① Handle multiple requests simultaneously
- ② Modify state of objects
- ③ Need for control of atomic actions
- ④ Regulate and schedule access to objects



# Active Object pattern

The Active Object pattern could do the trick, but...

Not always suitable because

- 1 More complicated then necessary
- 2 Request and execution both in separate thread
- 3 The ability of decoupling synchronization is not always needed
- 4 Scheduling and registration (of Activation List) only slows it down
- 5 Doesn't locate the synchronization closely to functionality



# Monitor Object pattern

Monitor Object pattern, Thread-safe Passive Object or  
Code Locking pattern

## Design

- 1 Create an object containing the unsafe implementation  
(Synchronized Method)



# Monitor Object pattern

Monitor Object pattern, Thread-safe Passive Object or Code Locking pattern

## Design

- 1 Create an object containing the unsafe implementation (Synchronized Method)
- 2 Create an object responsible for synchronization (Monitor Object)



# Monitor Object pattern

Monitor Object pattern, Thread-safe Passive Object or Code Locking pattern

## Design

- 1 Create an object containing the unsafe implementation (Synchronized Method)
- 2 Create an object responsible for synchronization (Monitor Object)
- 3 Assign a Lock object to the Monitor Object (Monitor Lock)



# Monitor Object pattern

Monitor Object pattern, Thread-safe Passive Object or Code Locking pattern

## Design

- 1 Create an object containing the unsafe implementation (Synchronized Method)
- 2 Create an object responsible for synchronization (Monitor Object)
- 3 Assign a Lock object to the Monitor Object (Monitor Lock)
- 4 Provide the implementation means to suspend and resume (Monitor Condition)





# Monitor Object pattern

Monitor Object pattern, Thread-safe Passive Object or Code Locking pattern

## Design

- 1 Create an object containing the unsafe implementation (Synchronized Method)
- 2 Create an object responsible for synchronization (Monitor Object)
- 3 Assign a Lock object to the Monitor Object (Monitor Lock)
- 4 Provide the implementation means to suspend and resume (Monitor Condition)
- 5 N clients can invoke now simultaneously



## Means provided by the pattern

- 1 Only one synchronized method at a time running within an object
- 2 Separation of the low level lock acquirement and the synchronized implementation
- 3 Ability to suspend and resume execution within an method
- 4 Tightly coupled locking mechanism for increased performance



# 1: Free toilet




## 2: Toilet taken and locked



Separation of toilet and lock



### 3: Next people waiting in order of arrival




## 4: People unlock and leave



Enabling next waiting to go

## 5: People who can't deliver yet



Don't hold the lock on constipation



## 6: People want to continue




Just continue when you are feeling ready (notified by digestion)





## 7: Multiple objects with each their own lock

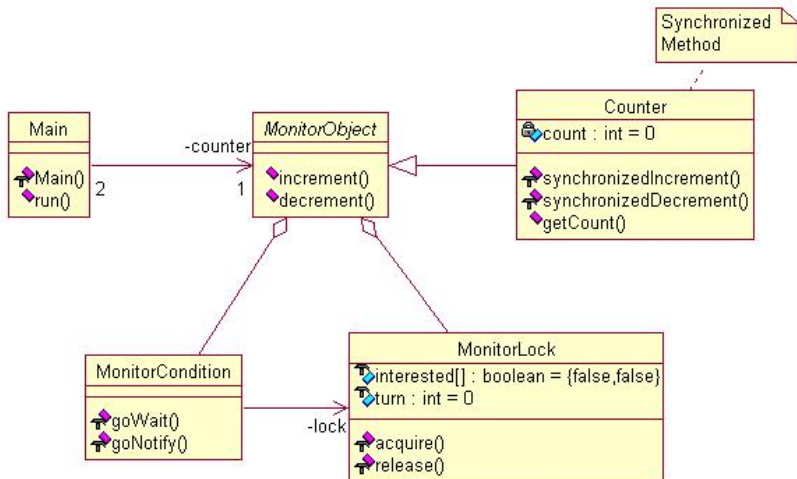


# Pitfalls

- 1 Not calling wait() in stable state  
Pull your pants up first !
- 2 Forget to release lock when an exception occurs  
Children panic when they can't open the door
- 3 Not making the method synchronized when needed  
Don't forget to lock your toilet door !



# Example: synchronized implementation (1)



## Example: synchronized implementation (2)

```
// Synchronized Method
class Counter extends MonitorObject
{
    private int count = 0;

    void synchronizedIncrement(int threadId)
    {
        //wait(threadId);    // will hold after a while, because

        count++;
        System.out.println("Thread " + threadId + ": " + count);
    }

    void synchronizedDecrement(int threadId)
    {
        //notify(threadId); // not every notify has an increment
        // waiting

        int current = count;
        current--;    // buffer with latency
        try{ Thread.currentThread().sleep(10); } catch(Exception
        count = current;
        System.out.println("Thread " + threadId + ": " + count);
    }

    public int getCount(){ return count; }
```



## Example: synchronized implementation (3)

```

public abstract class MonitorObject
{
    private MonitorLock    lock    = new MonitorLock();
    private MonitorCondition condition = new MonitorCondition(lock);

    public void increment(int threadId)
    {
        lock.acquire(threadId);

        synchronizedIncrement(threadId);

        lock.release(threadId);
    }

    public void decrement(int threadId)
    {
        lock.acquire(threadId);
        synchronizedDecrement(threadId);
        lock.release(threadId);
    }

    void wait  (int threadId) { condition.goWait(threadId); }
    void notify(int threadId) { condition.goNotify(threadId); }

    abstract void synchronizedIncrement(int threadId);
    abstract void synchronizedDecrement(int threadId);
    public abstract int getCount();
}

```



# Example: synchronized implementation (4)

```
class MonitorCondition implements Runnable
{
    private MonitorLock lock;
    private volatile int threadId;

    MonitorCondition(MonitorLock lock)
    {
        this.lock = lock;
    }

    void goWait(int threadId)
    {
        lock.release(threadId);

        synchronized(this)
        {
            try{ wait(); } catch(Exception e){}
        }
    }

    void goNotify(int threadId)
    {
        this.threadId = threadId; // not thread safe, so
        new Thread(this).start(); // only suited for 2 threads
                                // non-blocking request of lock
                                // (dead-lock otherwise)
    }

    public void run()
    {
        lock.acquire(1 - threadId);
        synchronized(this){ notify(); }
    }
}
```



# Example: synchronized implementation (5)

```
// Dekker (1965)
class MonitorLock
{
    // both not yet interested in the lock
    volatile boolean[] interested = { false, false };
    volatile int turn = 0;        // may also be 1

    void acquire(int id)          // thread id = 0 or 1
    {
        int other = 1 - id;
        interested[id] = true;    // show interest of taking the lock

        while(interested[other]) // while the other one also is interested
        {
            // if the other thread has got the turn
            if(turn == other)
            {
                // give up interest as long as the other has its turn:
                interested[id] = false;
                while(turn == other)
                    continue; // busy waiting

                // now he has lost the turn, say you're interested
                interested[id] = true;
            }
        }
    }

    void release(int id)
    {
        interested[id] = false;    // finished, so not interested in the lock
        turn = 1 - id;            // give the other thread the turn
    }
}
```



# Sharing a lock among several objects

- 1 Lock only when executing a part within a certain method
- 2 Expand the locking to affect several objects instead of just one





# Locking less than a whole method

In the counter example

```
void synchronizedDecrement(int threadId)
{
    //...

    lock.acquire(threadId);

    int current = count;
    current--;
    try{ Thread.currentThread().sleep(10); } catch(Exception e){}
    count = current;

    lock.release(threadId);

    // ...
}
```

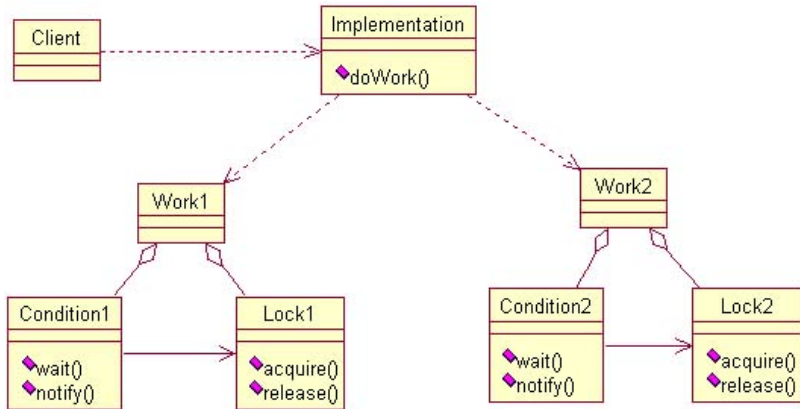
Ugly and losing abstraction

The Java way:

```
void partly()
{
    synchronized(this)
    {
        //...
    }
}
```



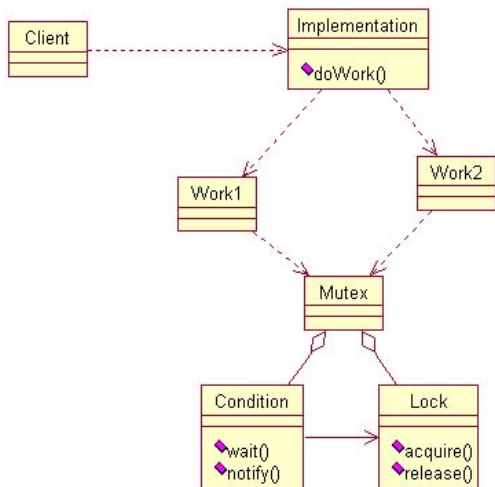
# Each object it's own lock



Worries of deadlock



# All objects sharing lock



Mutex provides extra influences



# A mutex using the Monitor Object pattern

```
public class Mutex
{
    //acquired == true when this Mutex is 'given away' to one thread
    volatile boolean acquired = false;
    Thread thread = null;

    public synchronized void acquire()
    {
        // while a thread has to wait
        while(acquired)
        {
            try
            {
                // if the thread doesn't already have the access
                if(this.thread != Thread.currentThread())
                    wait(); // let him wait
            }
            else
                break;
        }
        catch(Exception e){}

        // let the other threads wait
        this.thread = Thread.currentThread();
        acquired = true;
    }

    //...
}
```

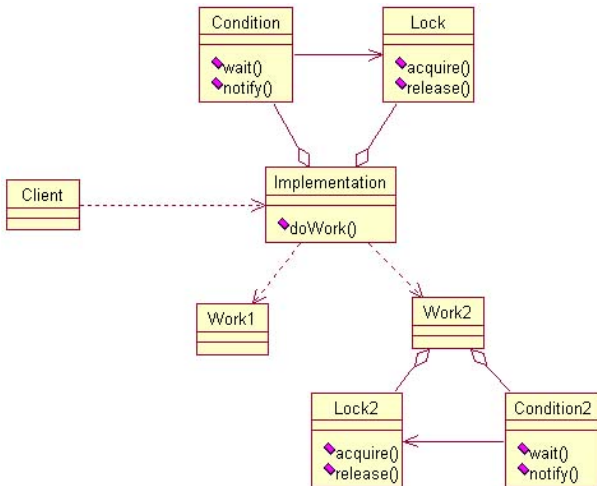


# A mutex using the Monitor Object pattern cont'd

```
//...  
  
public synchronized void release()  
{  
    // only the thread that called the acquire can  
    // release the lock  
    if (acquired && this.thread == Thread.currentThread())  
    {  
        // wake other waiting threads up  
        thread = null;  
        acquired = false;  
        notify();  
    }  
}
```



# The pattern doesn't always simplify it



# New problems created

- ① Nested monitor acquisitions
- ② Inherited methods not automatically synchronized
- ③ Starvation while it shouldn't
- ④ Deadlocks



## Deadlock avoidance

- 1 Don't call synchronized methods in synchronized bodies when not strictly necessary





## Deadlock avoidance

- 1 Don't call synchronized methods in synchronized bodies when not strictly necessary
- 2 Using timeout mechanisms releasing held locks combined with exceptions  
(dangerous for timeouts also can occur due to e.g. network failure)



## Deadlock avoidance

- 1 Don't call synchronized methods in synchronized bodies when not strictly necessary
- 2 Using timeout mechanisms releasing held locks combined with exceptions  
(dangerous for timeouts also can occur due to e.g. network failure)
- 3 Only use one lock for entire system  
(only usefull if performance really isn't a matter)



## Deadlock avoidance

- 1 Don't call synchronized methods in synchronized bodies when not strictly necessary
- 2 Using timeout mechanisms releasing held locks combined with exceptions  
(dangerous for timeouts also can occur due to e.g. network failure)
- 3 Only use one lock for entire system  
(only usefull if performance really isn't a matter)
- 4 Always acquire locks in the same order  
(not very reliable though !)



## Deadlock avoidance

- ① Don't call synchronized methods in synchronized bodies when not strictly necessary
- ② Using timeout mechanisms releasing held locks combined with exceptions  
(dangerous for timeouts also can occur due to e.g. network failure)
- ③ Only use one lock for entire system  
(only usefull if performance really isn't a matter)
- ④ Always acquire locks in the same order  
(not very reliable though !)
- ⑤ Keep registration of:  
Which locks currently acquired by which method  
Which locks needed for the method to execute  
(and then implement `canThreadWaitOnLock()` )



# High level use of mutexes and the Monitor Object pattern

## `java.util.concurrent.locks`

Provides extended capabilities (introduced in Java 1.5)

- 1 Non-blocking attempt to acquire a lock using `tryLock()`
- 2 2 acquires from same thread needs 2 releases
- 3 Provides information like `'isLocked'` and `'getLockQueueLength'`
- 4 Provides means to extend the amount and types of conditions to wait on
- 5 Ease use of timeout mechanisms



# Performance penalty

## Performance in Java

### Synchronized vs Non-synchronized method

Incrementing and decrementing value 200 million times in a loop with only one thread

JDK version	Synchronized	Not Synchronized	$\delta$ in ms
1.1.8	1032 ms	1016 ms	16
1.4.2	1859 ms	1421 ms	438
1.5.0	2141 ms	1719 ms	422



# Performance penalty

Monitor object principle used by:  
Hashtable, Vector, StringBuffer, util.Properties

## Hashtable vs HashMap

Tested on Maps containing 4.000 entries  
and average based on 500 runs

	Hashtable	HashMap
fill	2 ms	3 ms
iterate	154 ms	111 ms
remove	6 ms	5 ms



# Performance penalty

## Performance in Java

### Synchronized vs Non-synchronized method

Calling `System.getProperties()`, hard for the VM to optimize,  
500 times using 10 threads

	<code>void method()</code>	<code>synchronized(this){...}</code>	<code>synchronized method()</code>
avarage	62 ms	80 ms	94 ms
maximum	186 ms	240 ms	282 ms





# The Pros and Cons of the pattern

## Benefits

- ① Simplification of concurrency control implementation
- ② Simplification of scheduling method execution
- ③ Implementation 'separated' from concurrency control
- ④ Locking mechanism and implementation closely coupled



# The Pros and Cons of the pattern

## Drawbacks

- ① Locking mechanism and implementation closely coupled
- ② Limited amount of control (e.g. no reordering of calls)
- ③ Limited scalable
- ④ Complicated extensibility semantics
- ⑤ Inheritance anomaly
- ⑥ Nested monitor lockout

## Unavoidable drawbacks

- ① Concurrency remains complicated
- ② Caching at different levels
- ③ Almost impossible to test
- ④ Big responsibility on the programmer just implementing methods



# Questions ?



# Exercise

In a situation where there are several layers of caching, can you think up of an structure how you would design the Monitor Object pattern within that situation ? Which object(s) would have the lock, what objects/methods could get called simultaneously, who would invoke who, how many locks would you have and why ?

Draw some UML diagram

Briefly explain your intention

Some help:

```
class A
int value = b.get();
b.set(value-1);
```

```
class B
int get()
{
    int value = c.get();
    c.set(value-1);
    return value;
}

void set(int value){
    c.set(value); }
```

